

# ccnSim user manual

October 17, 2014



# Contents

<b>1</b>	<b>ccnSim overview</b>	<b>1</b>
1.1	What is ccnSim? . . . . .	1
1.2	ccnSim and Oment++ . . . . .	1
1.3	Overall structure of ccnSim . . . . .	2
1.4	Downloading and installing ccnSim . . . . .	3
1.5	Run your first simulation . . . . .	4
1.6	Organization of this manual . . . . .	4
<b>2</b>	<b>The ccnSim simulator</b>	<b>5</b>
2.1	Topology definition . . . . .	5
2.2	Content handling . . . . .	6
2.3	Nodes, Content Store, and Strategy Layers . . . . .	7
	2.3.1 Core Layer . . . . .	7
	2.3.2 Caching strategies . . . . .	7
2.4	Clients . . . . .	10
2.5	Output . . . . .	10
	2.5.1 Output files . . . . .	10
	2.5.2 Statistics . . . . .	11
2.6	Summary . . . . .	12
<b>3</b>	<b>Practical ccnSim</b>	<b>15</b>
3.1	Run a simple simulation . . . . .	15
3.2	Using the graphical interface . . . . .	19
3.3	Simulating ranges . . . . .	20
3.4	Summary . . . . .	21



# Chapter 1

## ccnSim overview

### 1.1 What is ccnSim?

ccnSim is a scalable chunk-level simulator for Content Centric Networks (CCN)[6], that we developed in the context of the Connect ANR Project.

- It is written in C++ under the Omnet++ framework.
- It allows to simulate CCN networks in scenarios with large orders of magnitude. In particular, ccnSim is able to simulate content stores up to  $10^6$  chunks and catalog sizes up to  $10^8$  files in a reasonable time.
- It is distributed as free software, downloadable at <http://www.enst.fr/~drossi/ccnsim>

Roughly speaking, ccnSim extends Omnet++ as to provide a modular environment in order to simulate CCN networks. Mainly, ccnSim models the forwarding aspects of a CCN network, namely the caching strategies, and the forwarding layer of a CCN node. However, it is fairly modular, and simple to extend. We hope that you enjoy ccnSim in which case we ask you to please cite our paper [5].

### 1.2 ccnSim and Oment++

Omnet++ is a C++ based event-driven framework. It provides: i) a set of core C++ classes, to extend for designing the behaviour of a custom simulator; ii) a network description language (**ned**) to describe how the custom modules interact each other; iii) a **msg** language defining the messages node will exchange.

Our simulator comes as a set of custom modules and classes that extend the Omnet++ core in order to simulate a CCN network. A ccnSim simulation steps traverses three phases:

- Compiling ccnSim source files and linking with the Omnet++ core.

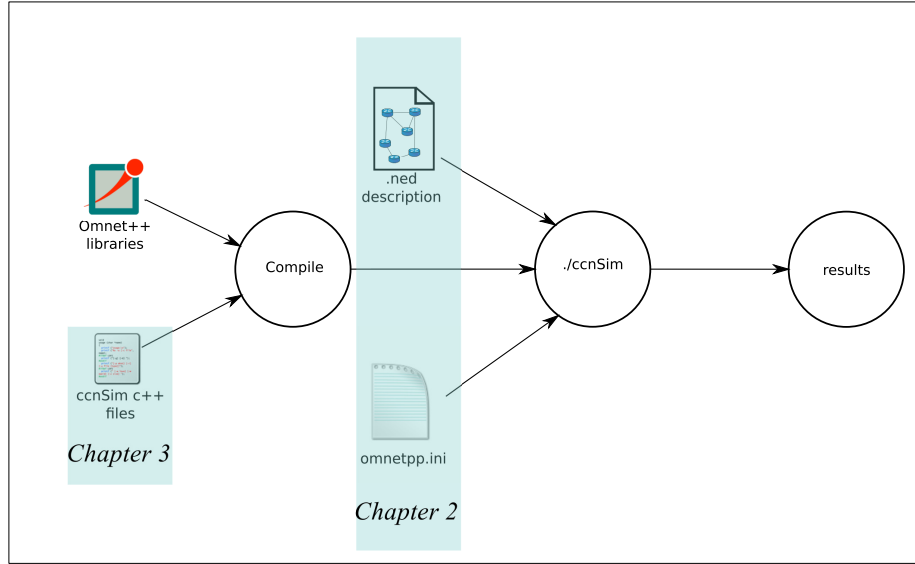


Figure 1.1: Basic steps in a `ccnSim` simulation.

- Writing the description of the topology (usually the user will need only to set up connections between the CCN nodes).
- Initializing the parameters of each module. This can be done either directly from the `ned` files or from the `omnetpp.ini` initialization file.
- Executing the simulation.

We report the aforementioned steps in Fig. 1.1. In the remainder of this manual we assume the reader has a basic knowledge of the Omnet++ environment. Otherwise we invite the interested reader to give a look at [1].

### 1.3 Overall structure of `ccnSim`

In order to better understand `ccnSim` structure, we provide a look at its internal directory structure:

As said within the introduction, `ccnSim` is a package built over the top of Omnet++. As such, we can divide its structure in three different subunits.

**Module properties** are defined in its `ned` file and simply specify the module appearance. . For instance, for a cache we could specify its size, or its replacement algorithm. These files are all contained within the `modules` directory.

```

|-- topologies
|-- modules
|   |-- clients
|   |-- content
|   |-- node
|       |-- cache
|       |-- strategy
|   |-- statistics
|-- packets
|-- include
|-- src
|   |-- clients
|   |-- content
|   |-- node
|       |-- cache
|       |-- strategy
|   |-- statistics

```

**Module behaviours** represent ccnSim core and are defined as C++ classes contained within the **src** and **include** directories (sources and headers, respectively). Roughly, for each module in the **modules** directory a C++ class is defined specifying its behaviour.

**Packets** Data and interest packets of the CCN architecture are defined within the **packet** directory. Messages are specified in the **msg** syntax.

Finally, the **topologies** directory contains a set of sample topology that the user can employ for starting with ccnSim (an example with Abilene network will be provided in Chap. 3).

## 1.4 Downloading and installing ccnSim

You can freely download ccnSim from the ccnSim home page <http://www.enst.fr/~drossi/ccnsim>.

If you do not have Omnet++ (version  $\geq 4.1$ ) installed, download it from <http://www.omnetpp.org> and install it following the “InstallGuide” of Omnet++, that you usually find in the documentation section of Omnet website.

The new version of ccnSim makes use of the boost libraries[2], thus you should have a minimal boost installation on your system. To install boost, download it from <http://www.boost.org/users/download/> and extract the content of the archive. Inside the parent directory, you will find a directory named **boost**. Move it inside **/usr/include/**. You may need superuser privileges.

In order to install `ccnSim`, it is first necessary to patch `Omnet++`. Then, you can compile the `ccnSim` sources. These steps are as follows:

```
john:~$ cd CCNSIM_DIR
john:CCNSIM_DIR$ cp ./patch/ctopology.h OMNET_DIR/include/
john:CCNSIM_DIR$ cp ./patch/ctopology.cc OMNET_DIR/src/sim
john:CCNSIM_DIR$ cd OMNET_DIR && make && cd CCNSIM
john:CCNSIM_DIR$ ./scripts/makemake
john:CCNSIM_DIR$ make
```

In this snippet of code, and for the remainder of the document, we suppose that `CCNSIM_DIR` and `OMNET_DIR` contain the installation directory of `ccnSim` and `Omnet++` respectively.

## 1.5 Run your first simulation

To run your first simulation

```
john:CCNSIM_DIR$ ./ccnSim -u Cmdenv
```

If you have kept the `omnet.ini` files that is by default in `ccnSim` package, at the end of the simulations you should find different output files for different values of `alpha` and different decision policies. To see the files:

```
john:CCNSIM_DIR$ cat results/abilene/\
F-spr/D-{decision}/R-lru/alpha-{alpha_value}/ccn-id0.sca
```

where you should replace `{alpha_value}` with 0.5, 0.6, 0.7, 0.8, 0.9 or 1 and `{decision}` with `fix0.1`, `lcd`, `lce`, `prob.cache`.

For more details about launching simulations, please refer to Chap. 3.

## 1.6 Organization of this manual

This document is organized as follows:

- In Chap. 2 we give a description of the module organization of `ccnSim` together with a brief description of the most important parameters that describe the simulation.
- Chap. 3 reports a brief `ccnSim` tutorial. We will show how to simulate and extending `ccnSim`.



## Chapter 2

# The ccnSim simulator

In the following we give an Omnet++ perspective of ccnSim. Note that the following parameters can be set both from the `.ini` and from within the `.ned` description file of each module[1]. Most of the modules illustrated have a corresponding C++ class.

If you start to look at the source code, to deeply understand how ccnSim works, you will find some blocks of code inside the following preprocessor directives

```
#ifdef SEVERE_DEBUG
... code ...
#endif
```

The nested lines of code are disabled by default. You can skip reading them as they have no functional role and will not help you understand ccnSim. They serve only to debug ccnSim. It is highly advisable to enable the `SEVERE_DEBUG` mode when you modify ccnSim source code, in order to check if you are introducing inconsistencies or if you are experiencing strange errors which are easier to debug in the `SEVERE_DEBUG` mode. For more details, search for `SEVERE_DEBUG` in `include/ccnsim.h`.

### 2.1 Topology definition

The `network` represents the top level module of a ccnSim simulation. In there, the user should define the connections between different CCN `nodes` modules and the placement of clients and repositories as well as the number of nodes within the network used. Each network module, MUST extend the `base_network`, in which the other modules (i.e., clients, statistics, and so forth) are defined.

**Clients placement** Clients represent an aggregate of users: thus, at most one client is connected and active on a given node. Indeed, in ccnSim clients are connected to each node of the network. You can see in `networks/base.ned`

that port 0 of each node is connected to a client. The placement consists in specifying how many (and which) of them are active. The activation of client takes place in the method `client::initialize()`. The basic parameters for client placement are defined in `networks/base.ned`:

- **num\_clients**: this integer value specifies how many clients are active over the network.
- **node\_clients**: comma separated string that specifies which CCN node has an active client connected to itself. The number of clients specified should be  $\leq$  **number\_clients**. If the number of clients specified is  $<$  **number\_clients** (this includes the case of an empty string ' ') the remaining clients are distributed randomly across the network.

To understand which node the  $i$ -th active client is attached to, open the output `.sca` file (see Sec. 2.5.1) and search for `client-i`.

**Repositories initialization** In `ccnSim` there is no real node representing a repository. A CCN node just *knows* that he owns a repository connected to itself. The repository distribution takes place in `content_distribution::init_repos()`. The distribution of repositories basically depends on two parameters, defined in `networks/base.ned`:

- **num\_repos**: integer value that specifies how many repositories should be distributed over the network.
- **nodes\_repos**: comma separated string that specifies which CCN node has a repository connected to itself. The number of repositories specified should be at most **number\_repos**. If the number of repositories specified is  $\leq$  **number\_repos** (this includes the case of an empty string ' ') the remaining repositories are distributed randomly across the network.

To understand which node the  $i$ -th repository is attached to, open the output `.sca` file (see Sec. 2.5.1) and search for `repo-i`.

## 2.2 Content handling

The `content_distribution` module takes no part in the architecture itself, but accomplishes many crucial tasks for the correct functioning of the simulation.

**Catalog initialization** The *catalog* is a table of contents. We will use interchangeably the terms *content* or *object*. The catalog is described by these parameters.

- **objects**: represents the cardinality of the catalog, expressed in number of contents.

- **file\_size**: as the contents are distributed like a geometric distribution, this parameter represents the average size of the file, in chunks. Moreover, if **file\_size** is set to one, whole objects are considered (each one composed by a single chunk).
- **replicas**: this parameters indicates the degree of replication of each content, and has to be less  $< \text{num\_repos}$ . In other words, the  $i$ -th content will be (randomly) replicated over exactly **replicas** repositories.

**CDF initialization** For the time being, the only distribution implemented is a Mandelbrodt-Zipf. The **content\_distribution** module (corresponding to the **content\_distribution** class takes, besides the number of objects (**objects** parameter), other two parameters: **alpha** is the shaping factor of the MZipf, and **q** represents the MZipf plateau.

## 2.3 Nodes, Content Store, and Strategy Layers

The **node** module represents the ccnSim core. It is compounded by three other submodules: the **core\_layer**, the **strategy\_layer**, and the **content\_store** modules. Each of these submodules are described below jointly with the parameters they accept.

### 2.3.1 Core Layer

The **core\_layer** module implements the basic tasks of a CCN node, and the communication with the other node's submodules. It handles the PIT, sending data toward the interested interfaces. It handles the incoming interests by sending back data (in the case of a *cache hit* within in the **content\_store**), or by appending the interest to the existent PIT entry. In the case no entry exists yet it queries the **strategy\_layer** in order to get the correct output interface(s).

### 2.3.2 Caching strategies

We can think to a caching algorithm on a CCN network as a triple  $\langle \mathcal{F}, \mathcal{D}, \mathcal{R} \rangle$ . The forwarding strategy  $\mathcal{F}$  determines which path exploiting, i.e., where the given Interest has to be sent. The decision policy  $\mathcal{D}$  returns a boolean value saying if the current node on the path has to cache or not the given data. The replacement  $\mathcal{R}$  drops an element from a full cache to make room for the incoming (cacheable) element. In the following, we explain how to set each element of this triple within ccnSim. We recall that these are pa

#### Forwarding strategies - FS

The forwarding strategy receives an interest for which no PIT entry exists yet. Then, it decides on which output face the interest should be sent. We make the assumption that each node knows the location of the permanent copy of each

content. There are different strategies actually implemented within `ccnSim`. One particular strategy can be chosen by setting the `FS` parameter of a node module.

**Shortest Path Routing - FS = `spr`** The strategy layer choses the shortest path repository and sends packets on the corresponding interface.

**Random Repository - FS = `random_repository`** The strategy layer choses one repository at random out of the given set of repository. Note that this strategy requires that the core nodes follow the path chosen by the edge node (the node to which the client is attached to).

**Nearest Replica Routing (Two phases) - FS = `nrr`** With this setting, the strategy layer first explores the neighboring nodes by flooding meta-interest (i.e., interests which do not change the content of the cache) with a given TTL. Then, the strategy sends the interest packet toward the nearest nodes having the content available. The TTL can be set by the means of the `TTL1` parameter. Setting `TTL1` to  $\infty$  (i.e., greater than the network diameter) makes NRR degenerating in the ideal NRR (`iNRR`) strategy, which explore the entire network, looking for a copy of the given content.

**Nearest Replica Routing (One phase) - FS = `nrr1`** In this last case, a node which receives an interest sets up an exploration phase, in which the node floods the neighborhood with the request for the given object. When the copy is found the data comes back (it can be a permanent copy, or a *cached copy* as well). The scope of the flooding can be set by the means of `TTL2` parameter.

Note that the routing table is learned at run-time by the mean of the Dijkstra algorithm. It is possible passing the routing matrix as an ASCII file, specified by the `file_routing` parameter of the `strategy_layer` module.

### Decision Strategies

By connecting more caches there arises the problem of caching coordination (who caches what?). Thus, a cache uses a decision strategy (or meta-caching algorithm) in order to decide if storing or not the incoming data. The parameter `CD` in `omnetpp.ini` sets which decision algorithm to employ for the given cache.

- `DS = lce` implements the Leave Copy Everywhere policy. Store each incoming chunk within the cache. It is defined in `include/always_policy.h`.
- `DS = lcd[7]` implements the Leave Copy Down policy. Store each incoming chunk only if is the downstream node of the (permanent or temporary) retrieved copy. It is defined in `include/lcd_policy.h`.
- `DS = btw[4]` implements the Betweenness Centrality policy. On a given path, only the node with highest betweenness centrality stores the chunk. It is defined in `include/betweenness centrality.h`.

- `DS = fixP[8]` implements the Fixed probability decision. The parameter `P` indicates the probability with which a given node stores the incoming chunk. Example `DS = fix0.1`. It is defined in `include/fix_policy.h`.
- `DS = prob.cache[9]` implements the ProbCache strategy. As much a node is far from the (either temporary or permanent) copy, the less is the probability of caching the given element. It is defined in `include/prob_cache_policy.h`.
- `DS = never` disable caching within the network. It is defined in `include/never_policy.h`.
- `DS = costawareP[3]` implements the Cost-Aware (CoA) strategy. An incoming object is accepted with a probability proportional to its price. `P` is the average cache admission probability weighted on the different prices. It is defined in `include/cost_related_decision_policies/costaware_policy.h`. See `cost_scenario.tar.gz` to start experimenting with it.
- `DS = ideal.blind[3]` implements the Ideal-Blind strategy. An incoming object is accepted only if its rank is greater than or equal to the rank of the eviction candidate, i.e. the least recently used object stored in the cache. It is defined in `include/cost_related_decision_policies/ideal_blind_policy.h`. See `cost_scenario.tar.gz` to start experimenting with it.
- `DS = ideal.costaware[3]` implements the Ideal-CoA strategy. An incoming object is accepted only if its price is greater than or equal to the one of the eviction candidate, i.e. the least recently used object stored in the cache. It is defined in `include/cost_related_decision_policies/ideal_costaware_policy.h`. See `cost_scenario.tar.gz` to start experimenting with it.

To understand how the decision policy is set up, see `base_cache::initialize()` in `src/node/cache/base_cache.ned`.

### Replacement strategies

Finally, within the `node` module, the user can choose which type of caching using. This choice is fulfilled by the means of the `RS` parameter of the node compound module. The following is a brief description of the algorithms currently available within `ccnSim`.

- `RS = lru_cache` implements an LRU replacement cache. It is the most used algorithm within the literature, and simply replaces the least recently used item stored within its cache.
- `RS = lfu_cache` implements an LFU replacement cache. By the means of counters an LFU cache may establish which is the least popular content, and deleting it when the cache is full.

- `RS = random_cache` implements a random replacement cache. When the cache is full the canonical behaviour of a random replacement is to choose at random an element to evict.
- `RS = two_cache` implements an extension of LFU and random replacement. It takes two random elements, and then evicts the *least* popular one.
- `RS = fifo_cache` implement a basic First In First Out replacement. The first element entered in the cache is the first to be evicted, once the cache is full.

## 2.4 Clients

As said above, a `client` represents an aggregate of users modeled as a Poisson process. In the current implementation, a client asks for files chunk by chunk (i.e., the *chunk window* `W` is fixed to one). The following parameters describe the client behaviour (implemented in the correspondent C++ class).

- `lambda` is the (total) arrival rate of the Poisson process. Of course, this parameter can be different for each client.
- `check_time` represents a timer, expressed in seconds. Client modules check the state of every download each `check_time` seconds.
- `RTT` Represents the Round Trip Time of the network (expressed in seconds). If the time needed for downloading a file exceeds this value, the client assumes that the download is expired and does something (usually resend the interest packet for the expired download).

## 2.5 Output

### 2.5.1 Output files

The output is collected in standard Omnet++ files. In particular, coarse grained, and per node statistics, are collected within the corresponding `.sca` vector file. Instead, fine grained and per content statistics are collected within the `.vec` file. In case you run a lot of simulations, `.vec` files may grow up to GB. In this case you may want to disable them by adding the following line in `omnetpp.ini`:

```
**vector-recording = false
```

### 2.5.2 Statistics

The way in which statistics are taken in `ccnSim` is rather complex. More in general, one of issues in taking statistics within a network of caches, is “when” starting to collect samples. One could start at time  $t = 0$ , taking into account the period of time in which the caches are still empty (*cold start*). Otherwise, we could wait for caches that fill up (*hot start*).

Moreover, when comparing simulations with models, often the system is supposed to be *stable*. In other words, statistics should be considered only after the transient phase of the system is vanished. Identifying the transient of the system is not a simple task. In `ccnSim` things work in the following way.

First, we wait for the nodes (or a subset of them) to be completely full. After that, we wait for the system to be stable. Stabilization happens when the variance of the *hit probability* of each node goes below a threshold, that is fixed by the parameter `variance_threshold` in `modules/statistics/statistics.ned`. This is implemented by sampling the hit probability of each node, and then calculating the variance of the samples collected. The parameters that affects the statistics calculations are:

- **ts**: the sampling time of the stabilization metric (i.e., the hit probability).
- **window**: the window of samples for which the variance is calculated.
- **partial\_n**: the set of nodes for which waiting for filling and stabilization.
- **steady**: how much simulation time you want to run your simulation after the stabilization has happened.

All the time variables here are expressed in seconds. For the sake of the example, let's suppose `window=60s` and `ts=0.1s`. That means: each 100ms a sample is collected. When 60s of samples are collected (i.e., 600 samples) the variance is calculated and tested against the threshold. The `partial_n` parameter is useful in case there are few clients and shortest path is used. In this case, some node could remain empty, and waiting for it would mean a infinite simulation. Besides the hit probability, the other statistics are handled per single module (e.g., per client or per CCN node).

- **p\_hit** ( $p_{hit} = \frac{n_{hit}}{n_{miss}+n_{hit}}$ ): it defines the probability of finding a content within the node.
- **hdistance**: represents the number of hops that an interest travels before hitting a copy of the requested chunk.
- **elapsed**: the total time for terminating a download of a file.
- **downloads**: the average number of downloads terminated by a given client.
- **interest** and **data**: the average number of Interest and Data messages (respectively) handled by a node of the network.

For most of these statistics it's possible to average on the different contents and/or on the different nodes within the network:

- Coarse grained statistics: we output one synthetic value averaged on every content and for every node (coarse grained statistics).
- Per node statistics: we output  $n$  curves (where  $n$  represents the number of nodes of the network), one for each node. Each curve is averaged on every content.
- Per content statistics: we output  $N$  curves (where  $N$  represents the number of contents in the catalog), one for each content. Each curve is averaged on all the nodes of the network.
- Fine grained statistics: we output  $Nn$  curves (fine grained statistics), one for each node and each different content of the system.

## 2.6 Summary

In this chapter we overviewed the parameters set that affect the behaviour of the main `ccnSim` modules. In Tab. 2.1 we briefly sum up what previously shown, indicating the meaning and the value got by the given parameter. We recall that these parameters can be set either within the corresponding `ned` module, or within the initialization file `ini`. To locate the `ned` module that owns a certain parameter `param`, you can do as follows

```
john:CCNSIM_DIR$ grep -r <param> modules networks
```



Syntax	Meaning	Values
<code>num_repos</code>	Number of the repositories	Int > 0
<code>node_repos</code>	Nodes that own a repository	Comma-separated string
<code>replicas</code>	Degree of replication for each content	Int <= <code>num_repos</code>
<code>num_clients</code>	Number of active clients	Int > 0
<code>node_clients</code>	Nodes which have an active attached	Int <= <code>num_clients</code>
<code>lambda</code>	Poisson arrival rate	Double > 0
<code>RTT</code>	Timeout to indicate a download failure (in seconds)	Double > 0
<code>check_time</code>	Timer indicating how often a download should be checked	Double > 0
<code>file_size</code>	Average file size (in chunks)	Int $\geq$ 1
<code>alpha</code>	MZipf shaping factor	Double $\geq$ 0
<code>q</code>	MZipf plateau	Double $\geq$ 0
<code>objects</code>	Catalog cardinality	Int > 0
<code>FS</code>	Forwarding strategy	String
<code>TTL1</code>	TTL for the NRR1 strategy	Int > 0
<code>TTL2</code>	TTL for the NRR strategy	Int > 0
<code>routing_file</code>	Routing matrix when fixed routing is used	String
<code>DS</code>	Decision strategy	String
<code>RS</code>	Replacement strategy	String
<code>C</code>	Cache size (chunks)	Int > 0
<code>window</code>	Stabilization window (in samples)	Int > 0
<code>ts</code>	Sampling time (in seconds)	Double > 0
<code>partial_n</code>	Num. of nodes to be full for starting statistic collections	Int > 0 or = -1(all nodes)
<code>steady</code>	Time to simulate after stabilization (in seconds)	Double > 0

Table 2.1: Summary of the main ccnSim parameters.



## Chapter 3

# Practical ccnSim

In this chapter, we will go through a complete simulation of a CCN. We simulate the following scenario:

- A general network of caches: Abilene.
- Clients are connected at each node.
- The client's arrival rate is  $\lambda = 1req/s$ . It is the same for every node in the network.
- The catalog is composed by 1000 objects, each one single sized.
- Object popularity is distributed like a Zipf with shaping factor  $\alpha = 1$ .

### 3.1 Run a simple simulation

The first step for running a simulation is to define the network. As mentioned in Chap. 2, connections, and number of nodes are easily defined within a `.ned` file. In Snippet 1 we report the whole `ned` file which describes the network. As mentioned in Chap. 2 `abilene_network` has to extend `base_network`, and then specifying number of nodes and connections. If we omit extending `base_network` the basic modules do not get initialized, and the simulations would not be executed. Finally, the package `network` represents the actual directory where Omnet looks for the `ned` file (in this case `CCNSIM_DIR/networks`).

The next step is to set the correct parameters within the initialization (`ini`) file. We set the network name to Abilene (the network defined in Snippet 1). The parameters are set following the indication in Tab. 2.1, and the final `ned` file is shown in Snippet 2. If the reader has already read Chap. 2, understanding the `ini` file should not be hard.

The `$` notation is used within the `ini` in order to define `ini` variables. `ini` variables are handy for dealing with ranges of values for the simulation parameters. Let's suppose that we want simulate different `alpha` values, let's

```

package networks;
network abilene_network extends base_network{
parameters:
    n = 11;
connections allowunconnected:
    node[1].face++ <--> { delay = 5.48ms; } <--> node[0].face++;
    node[10].face++ <--> { delay = 3.80ms; } <--> node[0].face++;
    node[10].face++ <--> { delay = 5.02ms; } <--> node[1].face++;
    node[10].face++ <--> { delay = 1.68ms; } <--> node[9].face++;
    [...]
}

```

Snippet 1: Ned file for the Abilene simulation.

say from 0.5 to 1 with a 0.1 steps. One choice could be running `ccnSim` different times, each time with a different value of `alpha`. Moreover, for each execution one has to specify different output files, in order to not override the results.

Instead, by the means of the `$` notation, we can easily write:

```

**alpha = ${a = 0.1...1 step 0.1}

```

In this way `Omnet++` will execute 10 different runs of the same scenario, each with a different  $\alpha$ . Being `a` a `ini` variable, it can be used for automatically specifying (the directory name of) the output files(`output-scalar-file` and `output-vector-file`). In this way each simulation run, will be hold on a different file. This notation has been used in Snippet 2 for: random generator seed (`rep`), shaping factor(`a`), forwarding(`F`), decision (`D`) and replacement strategy (`R`). Just recall that the name of the `ini` variable does not forcedly correspond to the `ccnSim` parameter.

Now, let's try to execute the simulation. The simpler method (the more complex one will be explained in a while) is to write the following.

```

john:CCNSIM_DIR$ ./ccnSim -u Cmdenv

```

The option `-u Cmdenv` is for executing `ccnSim` in console mode (without running the graphic engine). In Snippet 3 we report the result of the execution, skipping out uninteresting parts. As we can see, caches are being filled after 134 secs the simulation is initialized. Then, `ccnSim` waits for the variance of each node's `p_hit` to be under a given threshold. That happens after the second round, thus clearing out the old statistics. After about 1800 secs, the simulator stops, flushing on the console some of the synthetic values (saved in the `sca` file). If we go through the directory `CCNSIM_DIR/results/abilene/F-spr/D-lce/R-lru/alpha-1/` we would find three files with the same name `ccn-id0` and different extensions. As said in Chap. 2, the `sca` file contains the coarse and per node metrics. Let's

```

#General paramteters
[General]
network = networks.${net=abilene}_network
seed-set = ${rep = 0}
##### Repositories #####
**.node_repos = ""
**.num_repos = 1
**.replicas = 1
##### Clients #####
**.node_clients = ""
**.num_clients = 11
**.lambda = 1
**.RTT = 2
**.check_time = 0.1
##### Content Distribution #####
**.file_size = 1
**.alpha = ${a = 1}
**.objects = 10^4
##### Forwarding #####
**.FS = "${F = spr}"
**.TTL2 = 1000
**.TTL1= 1000
**.routing_file = ""
##### Caching #####
**.DS = "${ D = lce}"
**.RS = "${ R = lru }_cache"
**.C = 10^2
##### Statistics #####
**.window = 60
**.ts = 0.1
**.partial_n = -1
**.steady = 3600/2
#####
output-vector-file =
${resultdir}/${net}/F-${F}/D-${D}/R-${R}/a-${a}/ccn-id${rep}.vec

output-scalar-file =
${resultdir}/${net}/F-${F}/D-${D}/R-${R}/a-${a}/ccn-id${rep}.sca

```

Snippet 2: Initialization file for the Abilene simulation.

```

[. . . ]


Setting up network 'networks.abilene_network'...
Initializing...
Initializing Zipf...
Initialization ends...
Start content initialization...
Content initialized

Running simulation...
** Event #1    T=0    Elapsed: 0.000s (0m 00s)
    Speed:      ev/sec=0    simsec/sec=0    ev/simsec=0
    Messages:   created: 25    present: 25    in FES: 23
Caches filled at time 134.6
0] variance = 0.113934
1] variance = 0.0499831
2] variance = 0.0846741
[. . . ]
0] variance = 0.0110749
1] variance = 0.0064429
2] variance = 0.00657124
[. . . ]
** Event #372608    T=2054.6    Elapsed: 1.040s (0m 01s)
Speed:ev/sec=358277    simsec/sec=1975.58    ev/simsec=181.353
Messages:   created: 121187    present: 22    in FES: 22
[ . . . ] Simulation stopped with endSimulation().
Calling finish() at end of Run #0...
p_hit/cache: 0.161627
Distance/client: 2.20726
Time/client: 0.011430159267
Downloads/client: 1821.64

End.

```

Snippet 3: Execution of the simulation

Figure 3.1: Simulation result: `p_hit` per node.

say we would produce the `p_hit` for each node of the network. Moreover we want sort the output in ascending value. The following snippet of code, is apt to do that:

```
john:CCNSIM_DIR/results$ grep 'p_hit\[\'' [...] /ccn-id0.sca|
> awk '{print $4}' |
> sort -n > p_hit.data
```

The plot of the file `p_hit.data` is shown in Fig. 3.1. In pasting and copying the previous snippet, take care of substituting the `"[...]"` with the right directory `CCNSIM_DIR/results/abilene/F-spr/D-lce/R-lru/alpha-1/`.

## 3.2 Using the graphical interface

Sometimes you may need to verify if you wrote the topology in the correct way before launching the entire simulation campaign. Or you may need to see how single packets move inside the network. For this kind of preliminary checks, it may be useful to launch `ccnSim` with the Omnetpp graphical interface. To do so:

```
john:CCNSIM_DIR$ ./ccnSim
```

Please refer to Omnetpp manual to learn how to interact the graphical interface.

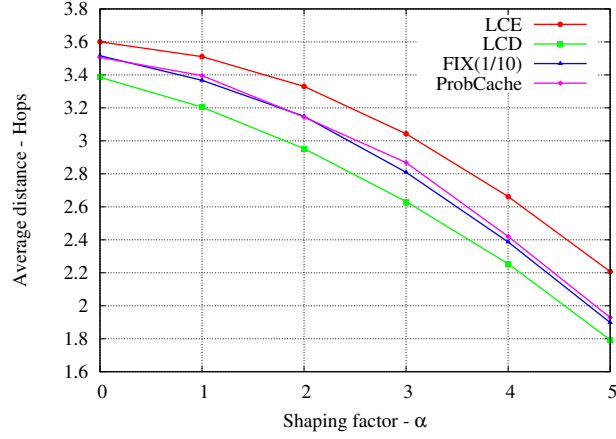


Figure 3.2: Comparison of different decision strategies for different values of  $\alpha$ .

### 3.3 Simulating ranges

As the last part of this brief `ccnSim` tutorial, we deal with simulating different decision strategies  $\mathcal{D}$  for  $\alpha \in [0, 1]$ . In particular, we are going to simulate `lce`, `lcd`, `prob_cache`, and `fix0.1`. In this last case, positive decision is taken one time over ten times, (otherwise stated, on a path long 10 hops only one node will cache the data). We will consider the average download distance `hdistance` as coarse grained metric. It turns out that, by exploiting the `$` notation, we can simulate this scenario in one single shot. First of all, we have to modify the `ini` file in the following way.

```
[ . . . ]
**alpha = ${a = 0.5...1 step 0.1}
[ . . . ]
**DS = ${D = lce,lcd,prob_cache,fix0.1}
[ . . . ]
```

As we see, we have a bunch of 24 simulations to run. Doing `./ccnSim -u Cmdenv` would execute the 24 runs in sequence. If we are running `ccnSim` on a machine more than a single processor, we would like to run a bunch of parallel simulations. The `runall.sh` script in the directory `CCNSIM_DIR/scripts`, launches the simulations in parallel on different processor. For modifying the number of processor, it suffices to open the file with a suitable text editor and



modify the option `-j` of the command. That said, by writing:

```
john:CCNSIM_DIR$ ./scripts/runall.sh
```

The whole set of simulations is executed in parallel.

Now, we have a bunch of 72 files. Indeed, for each different strategy we simulate 6 different values of  $\alpha$ , generating a `vec`, a `sca` and a `vci`. Each file is in the right directory with the name of the corresponding strategy and the corresponding  $\alpha$  value. Let's start evaluating the performance in terms of the (coarse grained) average download distance. As before some bash scripting does the magic:

```
john:CCNSIM_DIR/results$ grep 'distance [...] /ccn-id0.sca |  
> awk '{print $4}' > distance_lcd.data
```

Note that we are considering the LCD (`lcd`) decision strategy and that the string replacing the `"[...]"` should be `abilene/F-spr/D-lcd/R-lru/alpha-*` in order to get the right results. Indeed, for the given decision policy (in this case the LCD policy) we need all the different values for each different shaping factor. Repeating the above snippet for each strategy, and plotting the result, we generate Fig. 3.2.

## 3.4 Summary

In this section we have shown how to practically simulate a CCN network by the means of `ccnSim`. We simulated a simple scenario considering the Abilene network. Then we showed how to simulate a range of parameters, and producing the correct results.



# Bibliography

- [1] <http://www.omnetpp.org>.
- [2] Boost homepage. <http://www.boost.org/>.
- [3] A. Araldo, D. Rossi, and F. Martignon. Design and evaluation of cost-aware information centric routers. In *ACM ICN*, 2014.
- [4] W. Chai, D. He, I. Psaras, and G. Pavlou. Cache less for more in information-centric networks. In *IFIP NETWORKING*, volume 7289, pages 27–40. 2012.
- [5] R. Chiocchetti, D. Rossi, and G. Rossini. ccnSim: an highly scalable ccn simulator. In *IEEE ICC*, 2013.
- [6] Van Jacobson, Diana K. Smetters, James D. Thornton, Michael F. Plass, Nicholas H. Briggs, and Rebecca L. Braynard. Networking named content. In *ACM CoNEXT*, page 112, 2009.
- [7] N Laoutaris, H Che, and I Stavrakakis. The LCD interconnection of LRU caches and its analysis. *Performance Evaluation*, 63(7), 2006.
- [8] Nikolaos Laoutaris, Sofia Syntila, and Ioannis Stavrakakis. Meta Algorithms for Hierarchical Web Caches. In *IEEE ICPCC*, 2004.
- [9] I. Psaras, W. K. Chai, and G. Pavlou. Probabilistic in-network caching for information-centric networks. In *ACM ICN*, 2012.